

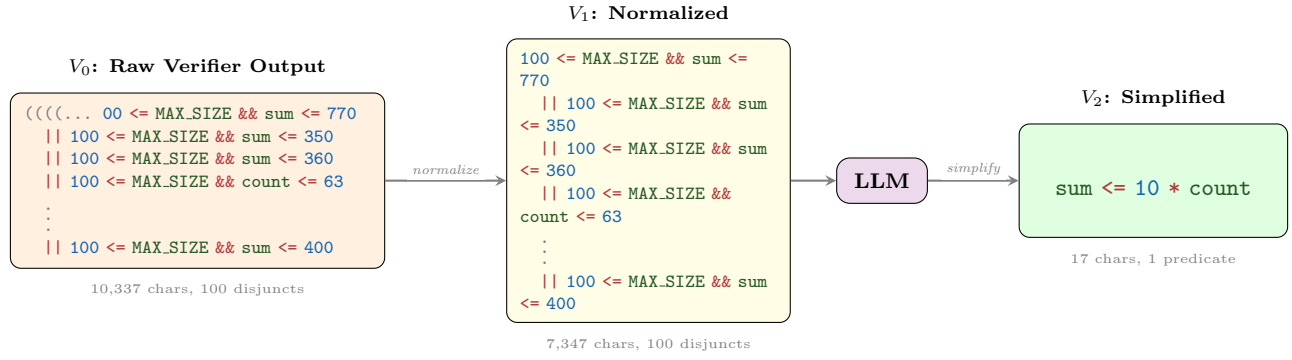


Figure 1: WONDA pipeline on 4820_1.c. The raw verifier output (V_0) is normalized (V_1); the LLM then tightens case-wise sum bounds to the linear inequality implied by each iteration (V_2): `sum <= 10 * count`, a 608× character reduction with **28.8**× verification speedup.