

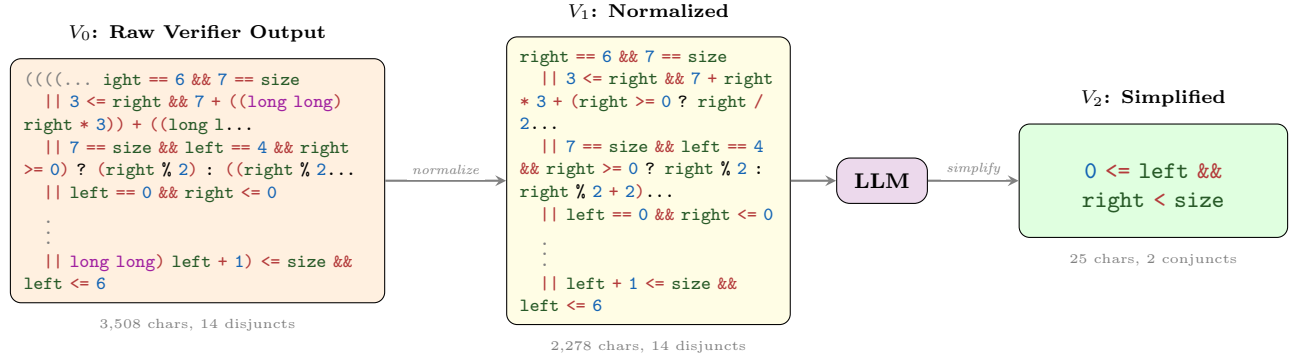


Figure 1: WONDA pipeline on 7861_1.c. The raw verifier output (V_0) is normalized (V_1); the LLM then removes size-specific overfitting and keeps valid binary-search index bounds (V_2): `0 <= left && right < size`, a 140× character reduction with **5.3**× verification speedup.